

Danyal Wahdat

mdanyalwahdat@gmail.com | +92 340 9731904 | linkedin.com/in/danyal-wahdat-b747a928b |
github.com/Danialpro2k04

PROFESSIONAL SUMMARY

AI LLMOps Engineer specializing in building scalable, production-grade orchestration layers. Experienced in architecting Enterprise RAG pipelines, semantic routing microservices, and deterministic agentic workflows using Python, FastAPI, and Qdrant. Passionate about bridging the gap between probabilistic LLMs and strict software engineering requirements by building resilient backend infrastructure.

TECHNICAL SKILLS

Languages & Frameworks: Python, Javascript, FastAPI, Bash/Shell Scripting

AI & LLMOps: LangChain, RAG Architecture, Semantic Routing, Agentic Workflows, Prompt Engineering

Cloud & Infrastructure: AWS, Azure AI Search, Docker, Apache Airflow, CI/CD pipelines

Databases & Tooling: Qdrant (Vector Databases), Git/GitHub, Linux, Cursor IDE

EXPERIENCE

Freelance Software Engineer

Remote

Upwork

Jan 2024 – Present

- Collaborated with international clients to deliver custom Python automation solutions, translating ambiguous business requirements into deployed, functional applications.
- Engineered an automated data extraction and web scraping pipeline using Python, efficiently extracting, cleaning, and structuring large datasets to save the client hours of manual data entry per week.

KEY PROJECTS

Multi-Agent AI Decision Debater (Flagship Project)

LangGraph, Groq, Tavily Search API

- Engineered a multi-agent cognitive architecture utilizing six specialized AI agents to challenge, verify, and synthesize complete analyses to eliminate costly "gut-trusting" in high-stakes decisions.
- Designed a complex 6-node graph topology featuring an Orchestrator for dimension planning, a parallel execution layer (Optimist, Devil's Advocate, Risk Analyst), a Researcher for live web exploration, and a Judge agent outputting verdicts with 1-5 harshness tuning and confidence scores.

Enterprise-Grade AI Search & RAG Backend

Azure AI Search, Vector DBs, Python

- Architected a production-ready Retrieval-Augmented Generation (RAG) system utilizing Azure AI Search to enable highly accurate, scalable, and secure document retrieval.
- Implemented custom context-retrieval mechanisms focusing on advanced text chunking and embedding strategies. Enforced strict prompt grounding logic, successfully eliminating LLM hallucinations in enterprise data environments.

Semantic Router Microservice

FastAPI, Qdrant, Python, Docker

- Built a high-speed API routing layer designed to intercept user intents locally before hitting external LLM endpoints. Utilized Qdrant vector embeddings to classify queries, bypassing expensive LLM calls for basic tasks, which optimizes operational API budgets and drastically reduces inference latency.

LLMOps Telemetry & Observability API

FastAPI, Observability, Python

- Engineered a robust backend telemetry system to monitor LLM performance across production AI applications. Developed endpoints to track live token usage, monitor API cost accumulation, and observe real-time inference latency to ensure system stability.

Predictive Maintenance Deep Learning Model

PyTorch, LSTM, Time-Series

- Created a foundational deep learning model to predict the Remaining Useful Life (RUL) of turbofan engines. Leveraged the complex NASA CMAPSS dataset to perform advanced time-series forecasting and data normalization.

EDUCATION

National University of Modern Languages (NUML)

Islamabad, PK

Bachelor of Science in Computer Science

Expected Graduation: 2027

Relevant Coursework: Operating Systems, Discrete Structures, Analysis of Algorithms, Database Systems, Object-Oriented Programming

CERTIFICATIONS

- **Large Language Model Operations (LLMOps)** - Duke University
- **Datacom - Automation AI Accelerator: From Co-pilot to Autonomous Agent Job Simulation** - Forage
- **Use AI as a Creative or Expert Partner** - Google
- **Generative AI Engineering and Fine-Tuning Transformers** - IBM
- **Vector Databases for RAG: An Introduction** - IBM
- **Google Prompting Essentials Specialization** - Google
- **Fundamentals of AI Agents Using RAG and LangChain** - IBM
- **Neural Networks & Deep Learning** - DeepLearning.AI